Splay Trees

Johan Montelius

August 25, 2017

Introduction

A splay tree is an ordered binary tree with the advantage that the last key we looked for is found in the root of the tree. We will rearrange the tree in every access, moving the key to the top and trying to keep the rest of the tree balanced.

The amortized cost of operations (search, insert, delete) are all $O(lg(n))$. There are worst case scenarios, since there is no guarantees that the tree is balanced, but we accept this since the data structure has its advantages. Frequently used keys will be found higher up in the tree so it is very good to use in an application where we expect temporal locality i.e. if you have used one key it is very likely that you will use this key again with in a short period of time.

In this assignment you're going to learn how to implement a quite tricky algorithm using pattern matching. You should look up some tutorial on splay trees so that you have a basic understanding of the algorithm but we will explain the algorithm as we go along. We will first look at the general idea of these operations using a graphical representation before going in to how to implement it in Erlang.

1   The splay tree

In this implementation we will create a splay tree that keeps key-value pairs. The key-values will reside in all of the nodes besides the leafs which are empty branches. The tree is ordered with smaller values in left branches.

Doing search in this tree is of course trivial but we will do a search where we also change the structure of the tree. This is the key operation in all splay tree operations; updating, lookup and insertion are essentially the same algorithms since they all perform the same transformation.

In this description we will first look at the operation to update or insert a value given a key. If the key is present we will update the value, if not we will add the key-value pair. Once we understand how this is done it will be easy to implement the other operations.

The splay operation is slightly different depending on if we are in the root of the tree or doing a recursive operation further down in the tree. Before explaining the general rules let's look at the rules for the root.

Figure 1: Zig: splay operation of the root when the key (K) is found in left branch.

1.1 splay of the root

Assume that we want to update a value for a key in a tree. We will have two very easy special cases when the tree is empty or if the key is found in the root of the tree. In these cases we simply return a tree with one node or update the existing root with the new value.

The splay operation comes in when we find the key in one of the branches. We should then rearrange the tree in a way that moves the key to root of the tree. The operation is quite straight forward as seen in the graph shown in figure 1.

Note that the lefter-most sub-tree, A, of course contains keys that are less than the key (K) that we're looking for and it is safe to have it as the left branch of the new tree. The sub-tree B contains keys that are greater than the key we are looking for but smaller than the key of the root (R). Its position in the transformed tree is therefore sound. The same thing goes for the sub-tree C that must contain keys that are greater than the key of the root.

There is s corresponding splay operations when we find the key in the left branch. Write it down using the same naming scheme as in in figure 1. This will make the implementation step easier.

If the key is not found in the root of either sub-tree, the more general splay operation is applied on either the left or right sub-tree. The splay operation will transform the sub-tree and create a situation where we can apply the root operation.

1.2 splay further down the tree

In general the splay operation on a tree will result in a tree where the key that we're updating becomes the root of the tree. The operation looks very similar to the operation that we have seen for the root of the tree, the only difference is that we now look further down the tree before determining what to do. We first describe the general rules before describing the special cases.

We have four cases and we could call them left-left, left-right etc but for

Figure 2: Zig-Zig: splay operation when key is found in left-left branch.

historic reasons we call them: zig-zig, zig-zag, zag-zig and zag-zag. We only show the two first since the two other are mirror images of the two first.

When describing these rules we use the naming 'K' for the node where the key is found, 'P' for the parent node and 'G' for the grandparent. Sub-trees that are moved around are called 'A', 'B' etc and we write them in an order so that we know that all keys in for example sub-tree 'A' are smaller than keys in sub-tree 'B'. We will later use these names for variables in our implementation so let's try to be consistent.

zig-zig

The zig-zig rule is used if we find the key in the lefter-left node. In figure 2 we see how we move the node holding the key we're looking for (K) to the root and rearrange the parent (P), and grandparent (G), to form the new tree. Note the order of the sub-trees A, B, C and D. Make sure that you understand why it is safe to do the transformation of the tree and why the tree is still ordered.

zig-zag

The second rule covers the case where we find the key in the left-right node. The transformation is a little different but the aim is the same; move the key to the root and rearrange the sub-trees to keep the tree ordered. In figure 3 we see how the transformation is done.

1.3   zag-zig and zag-zag

The two described rules of course have their mirror rules when the key is found in either the right-right or right-left node. The idea is the same, move the found key to the root and the parent node one level down. The

Figure 3: Zig-Zag: splay operation when the key (K) is found in left-right branch.

grandparent node becomes the child of either the parent node or the key node.

You're strongly advised to draw the graphs that describes these zag-zig and zag-zag rules. Keep the naming of nodes: G for grandparent, P for parent and K for the node where we find the key. Sub-trees are called: A, B etc and are named in order.

These are the complicated rules, what we have left are the very simple rules for some base cases.

1.4 zig or zag or nil

We have several base cases that cover the situations where the tree is not very deep or when we find the key that we're looking for in one of the two first levels. In the same way as for the root we could have an empty tree or a tree where the key is in the root; these cases are trivial. If the key is found in the root of either sub-tree we do a transformation that is exactly the same transformation that we would do for the root.

If the sub-tree where the key should be found is empty we do the same operation as we would do if we found the key in the root of the sub-tree. An example of what this looks like is seen in figure 4.

The rule is equivalent to the zig rule that we used in the root of the tree. It of course has its corresponding zag versions where the key should go into the right branch.

2 The implementation

We will implement a function `update/3` that takes a splay tree, a key and a value and returns a transformed splay tree where the key has been updated

Figure 4: Zig: splay operation when key should go into left branch.

or inserted if it was not present. In both cases the key value pair is of course found in the root of the transformed tree.

A tree is represented by:

- `nil`: the empty tree

- `{node, Key, Value, Left, Right}`: a node with a key, value and a left and right branch

In the implementation we will use the same naming of nodes as in the graphical representations. We will augment names with `k` or `v` depending on if we refer to the key or value. A key value pair will for example be named `Rk` and `Rv` if it's the root node. Nodes that are only moved around will be called `A`, `B` etc. The idea is that we should keep the implementation as close to the graphical description as possible.

2.1 update of the root

Let's start with the base cases for updating the root. We either have a empty tree or a node where we are lucky enough to find the key in the root. Below is the skeleton code that you should use, fill in the blanks.

```
update(nil, Kk, Kv) ->
    {node, ..., ..., ..., ...};

update({node, Kk, _, A, B}, Kk, Kv) ->
    {node, ..., ..., ..., ...};
```

The two general cases, *Zig* and *Zag* are where we will do a call to a general *splay* operation. The function `splay/2` takes a tree and a key and will return a tuple {`splay, Kv, A, B`} where `Kv` is the value of the key (`na` if the key is not found) and `A` and `B` the left and right sub-trees. When we do an update of a value we're not interested in the found value of the key but we will use this in the coming operations.

If we know that `splay/2` works for all trees then we can use it to implement the zig and zag rules of the root.

```
update({node, Rk, Rv, Zig, C}, Kk, Kv) when Kk < Rk ->
    %% The general rule where we will do the Zig transformation
    {splay, _, A, B} = splay(Zig, K),
    {node, ..., ..., ..., {node, ..., ..., ..., ...}};

update({node, Rk, Rv, A, Zag}, Kk, Kv) when Kk >= Rk ->
    %% The general rule where we will do the Zag transformation
    {splay, _, B, C} = splay(Zag, K),
    {node, ..., ..., {node, ..., ..., ..., ...}, ...}.
```

When you fill in the blank spots in the code, use variable names that found in the graphical representations of the rules. This will help you to verify that you're actually doing the right thing. If we now only can implement the `splay/2` function we are done.

2.2   splay of the tree

As for the root we have two base cases when the tree is empty and when the key is found in the root. These should be quite easy to write down.

```
splay(nil, _) ->
    {splay, na, ..., ...};

splay({node, Kk, Kv, A, B}, Kk) ->
    {splay, ..., ..., ...};
```

The general zig-zag rules only work if we have three levels to work with. We therefore have two more special cases where the left or right sub-tree is empty. We should still return a splay-tuple and therefore need to construct two sub-trees. Make sure that you get this right, what should go into the empty spaces.

```
splay({node, Rk, Rv, nil, B}, Kk) when  Kk < Rk ->
    %% Should go left, but the left branch is empty.
    {splay, na, ..., {node, Rk, Rv, ..., ...}};

splay({node, Rk, Rv, A, nil}, Kk) when  Kk >= Rk ->
    %% Should go right, but the right branch is empty.
    {splay, na, {node, Rk, Rv, ..., ...}, ...};
```

One more special case is when the key is actually found as the root in the left or right sub-tree. Very similar to the case above but now we actually have a found a value.

```
splay({node, Rk, Rv, {node, Kk, Kv, A, B}, C}, Kk) ->
    %% Found to the left
    {splay, ..., ..., {node, ..., ..., ..., ...}};

splay({node, Rk, Rv, A, {node, Kk, Kv, B, C}}, Kk) ->
    %% Found to the right
    {splay, ..., {node, ..., ..., ..., ...}, ...};
```

When this is done we are ready for the general zig-zag rules. We here list the complete coding of the zig-zig rule and you should complete the

remaining rules. The zig-zig rule applies the splay operation on the left-left sub-tree and obtains the sub-trees `A` and `B` that can then be used to rearrange the tree. Note that the splay operation also provides the found value, `Kv`, and this value is returned in the splay tuple.

```
splay({node, Gk, Gv, {node, Pk, Pv, ZigZig, C} , D}, Kk)
          when ((Kk < Gk) and (Kk < Pk)) ->
    %% Going down left-left, this is the so called zig-zig case.
    {splay, Kv, A, B}  = splay(ZigZig, Kk),
    {splay, Kv, A, {node, Pk, Pv, B, {node, Gk, Gv, C, D}}};
```

The zig-zag rule is similar but now we have walked down the left-right branch. Look at the graphical representation of the rule in figure 3. Fill in the blanks, which sub-trees should go where?

```
splay({node, Gk, Gv, {node, Pk, Pv, A, ZigZag} , D}, Kk)
          when ((Kk < Gk) and (Kk >= Pk)) ->
    %%  Going down left-right, this is the so called zig-zag case.
    ... =  splay(ZigZag, Kk),
    {splay, Kv, {node, Pk, Pv, A, B}, {node, Gk, Gv, ..., ...}};
```

The remaining two rules are the mirror images of the first two rules. If you have not done so yet you should draw the images of the rules using the same naming strategy. There should be a one-to-one mapping from drawn rules to clauses in the code. When your done you have implemented the update function for splay trees.

```
splay({node, Gk, Gv, A, {node, Pk, Pv, ZagZig, D}}, Kk)
          when ((Kk >= Gk) and ...) ->
      :
    {splay, Kv, {node, Gk, Gv, A, B}, {node, Pk, Pv, C, D}};

splay({node, Gk, Gv, A, {node, Pk, Pv, B, ZagZag}}, Kk)
          when  ... ->
      :
    {splay, ..., ..., ...}.
```

This is it, you have implemented the update function of splay trees. Write some small test examples that updates a tree, try this:

```
test() ->
   Insert =  [{3,c}, {5,e}, {2,b}, {1,a}, {7,g}, {4,d} {5,ee}],
   Empty = nil,
   lists:foldl(fun({K, V}, T) -> update(T, K, V) end, Empty, Insert).
```

3   why this exercise

It's of course fun to implement an algorithm that you might have hear off but not fully understood. The aim of this exercise is however not the algorithm per see but how pattern matching and recursion can be used to implement a fairly complicated algorithm.

I hope that you see how the graphical descriptions maps to patterns and how easy it is to build the desired transformation. If you would have implemented this algorithm from scratch, you would probably have to do some thinking and experimenting before coming up with the solution that we have now. Looking at the solution it looks obvious but it takes some time to understand how to get the recursion right.

3.1   a mutable solution

The Erlang implementation that we have now does not not look like the solution you would write in Java or C++. In Erlang, as almost all functional programming languages, data structures are immutable and we construct a new tree in each operation. In a language where you can change data structures your would most likely change the structure of the given tree.

If you implement the update procedure and change the structure of the tree, you will most likely work with a linked tree where all nodes know their immediate parent. The algorithm would then first traverse down the tree to find the node with the given key, or construct one in a leaf if it is not present, and then *splay* this node towards the root using the parent pointers.

You can probably work out what the algorithm looks like: if my parent has a parent and I'm the left child of my parent and my parent is the left child of its parent, then ...

```
void splay( node *x ) {
  while( x->parent ) {
    if( !x->parent->parent ) {
      if( x->parent->left == x ) right_rotate( x->parent );
      else left_rotate( x->parent );
    } else if( x->parent->left == x &&
               x->parent->parent->left == x->parent ) {
      right_rotate( x->parent->parent );
      right_rotate( x->parent );
    } else if( x->parent->right == x &&
               x->parent->parent->right == x->parent ) {
      left_rotate( x->parent->parent );
      left_rotate( x->parent );
    } else if( x->parent->left == x &&
               x->parent->parent->right == x->parent ) {
      right_rotate( x->parent );
      left_rotate( x->parent );
    } else {
      left_rotate( x->parent );
      right_rotate( x->parent );
    }
  }
}
```

The actual rotation operations are of course pointer manipulations where the left, right and parents pointers are set to form the new structure of the tree. Updating the data structure in place is of course more efficient than

building a new structure in each transformations. Immutable data structures do however have their advantage when it comes to robustness and possibility to run things in parallel. What we look at now is however how expressive the different approaches are when coding a complex algorithm.

```
void right_rotate( node *x ) {
    node *y = x->left;
    if(y) {
      x->left = y->right;
      if( y->right ) y->right->parent = x;
      y->parent = x->parent;
    }
    if( !x->parent ) root = y;
    else if( x == x->parent->left ) x->parent->left = y;
    else x->parent->right = y;
    if(y) y->right = x;
    x->parent = y;
  }
```

3.2   records and named fields

Since we are looking at how easy it is to express an algorithm we can take the opportunity to rewrite the program using Erlang records. The record functionality will give us the ability to use named fields instead of explicitly writing down a whole tuple in a pattern.

Assume that we have a record definition for a node whith the fields: `key`, `value`, `left` and `right`. We could then rewrite `update/3` in the following way (we here also combine the three alternatives in one clause).

```
-record(node{key, value, left=nil, right=nil}).

update(nil, Key, Value) ->
    #node{key=Key, value=Value, left=nil, right=nil};

update(#node{key=Kk, value=Kv, left=Left, right=Right}, Key, Value) ->
  if
    Key == Kk ->
    #node{key=Key, value=Value, left=Left, right=Right};

  Key < Kk ->
    {splay, _, A, B} = splay(Left, Key),
    #node{key=Key, value=Value, left=A, right={node, Kk, Kv, B, Right}};

  true ->
    {splay, _, B, C} = splay(Right, Key),
    #node{key=Key, value=Value, left={node, Kk, Kv,  Left, B}, right=C}
end.
```

Using the record construct we do not have to memorize the order of the elements in the node structure. It makes it easier to update the code, for example if we find that we need another element in each node. Using

9

explicit patterns is often the easier way of handling things while the use of more complex data structures almost requires that you use the record functionality.

4   Summary

Pattern matching is a powerful technique to describe the rules of a transformation. Together with recursion even a complex algorithm becomes quite easy to implement.